

UNIX and C Programming (COMP1000)

Lecture 7: Shell Scripting

Updated: 29th July, 2015

Department of Computing
Curtin University

Copyright © 2015, Curtin University
CRICOS Provide Code: 00301J

Reading

For more information, see the weekly reading list on Blackboard.

- ▶ Unfortunately, Hanly and Koffman do not cover shell scripting.
- ▶ However, the `bash` manual page is extremely detailed:

```
[user@pc]$ man bash
```

- ▶ Type the slash (/) character, followed by a search term, to find specific information in the page.

Outline

The Shell

Script Files

Files and Processes

Control

Data

Exit statuses

Regular Expressions

Shell Languages

- ▶ Each different shell provides an interpreted “scripting” language.
- ▶ You can write “scripts” — mini-programs that are read and executed by the shell.
- ▶ We *could* do everything with C. However. . .
- ▶ Scripts are able to perform certain tasks much more easily.
- ▶ One line of a script may translate to many lines of C code.

Scripting vs. Programming

Scripting languages are usually:

- ▶ Interpreted, not compiled;
 - ▶ (Scripts are read and executed by an “interpreter” program.)
- ▶ Based on strings and string substitution;
- ▶ Intertwined with external programs or “commands”;
 - ▶ (Scripts use external programs as a C program uses functions.)
- ▶ Geared towards manipulating files and processes;
- ▶ Not suitable for complex mathematical problems;
- ▶ Not suitable for writing large-scale software.

(There are exceptions. Some languages do not fall neatly into one category.)

Bash

- ▶ Bash is the de facto standard shell on most UNIX systems.
- ▶ Bash was inspired by the older “Bourne shell”, and stands for “Bourne Again SHell”.
- ▶ Despite radical differences, C and Bash are distantly related via ALGOL (as are most modern languages).
- ▶ Most other shells bear strong resemblances to Bash anyway.

Bash scripts can use the following programming-like features:

- ▶ variables and command-line parameters;
- ▶ if and case statements;
- ▶ while and for loops.

However, these all look a little bit different from C.

Hello World

This script will print “Hello world!”:

```
echo Hello world\!
```

To explain:

- ▶ The script contains a single “echo” command, which just outputs its arguments.
- ▶ Normally, “!” is a special character.
- ▶ We’ve escaped it (removed its special meaning) using a “\” (backslash).

We can run this script as follows:

```
[user@pc]$ bash helloworld.sh
```

(Assuming we named the script “helloworld.sh”).

Hash Bang

- ▶ We could *also* run the script like a normal program:

```
[user@pc]$ ./helloworld.sh
```

... if we make a couple of tweaks *beforehand*.

- ▶ First, can specify the interpreter inside the script itself:

```
#!/bin/bash
echo Hello world\!
```

- ▶ The very first characters must be “#!” (or “hash-bang”).
- ▶ In this case, the script is run using /bin/bash (the full path to bash).

- ▶ Second, we need to make the script file executable:

```
[user@pc]$ chmod 755 helloworld.sh
```

(We’ll return to this later.)

Comments

- ▶ Comments start with a “#” (hash), and are single-line.

```
#!/bin/bash

# Now, I'm going to ask you that question once
# more. And if you say no, I'm going to shoot
# you through the head. Do you have any cheese
# at all?
echo No, I was deliberately wasting your time.
```

(Most scripting languages use “#” to indicate comments.)

- ▶ Commenting serves the same purpose as in C, Java, etc.
- ▶ Use them liberally! (But make sure they're relevant...)

Variables — Revision From Lecture 2

- ▶ Recall that the shell (bash) has variables.
- ▶ These are strings, and are created through assignment:

```
name=Sam
age=65 # Still a string!
address="34 Green Street"
# Quotes are required if you need spaces.
```

- ▶ Recall that you access variables with a \$ sign:

```
echo $name is $age, and lives at $address.
```

```
echo ${name} is ${age}, and lives at ${address}.
```

This works by substitution, like makefile variables and C preprocessor constants.

- ▶ Note: if you access a variable that doesn't exist, you will get an empty string (*not* an error!)

Example Script

A brief taste of scripting:

```
#!/bin/bash

for file in $*; do
    if [ -x $file ]; then
        ./$file &
    fi
done
```

- ▶ This script demonstrates the use of for-each loops, if statements, variables and parameters — some of the common components of scripts.
- ▶ We'll discuss all of these in more detail...

Users

First, some background on UNIX.

- ▶ On UNIX, you have a UID (user ID), a numeric code corresponding to your username.
 - ▶ At Curtin, your UID and username are probably equal.
 - ▶ In general, usernames can be non-numeric.
- ▶ Your UID is recorded against every file and directory you create.
- ▶ That is, UNIX remembers who “owns” what.
- ▶ A file's owner decides who else can access it (roughly).

Root

- ▶ Much of a UNIX system is owned by “root” — UID zero (a.k.a. the superuser, operator, administrator, BOFH¹).
- ▶ The root user is a role, not an actual person.
- ▶ The root user has unlimited permission to access, modify and delete anything and everything.
- ▶ This power is only occasionally necessary.
- ▶ Overused, it can be very dangerous; mistakes are easily made, with catastrophic consequences.
- ▶ Even administrators have their own personal non-admin accounts, to limit the risk.

¹See “Bastard Operator From Hell”.

Processes

- ▶ A running program is called a “process”.
- ▶ Processes have a PID (process ID).
- ▶ Processes also have an owner — a UID (user ID).
- ▶ Conceptually, every process runs *as a particular user* — normally the user who starts the process.

Process Listing

- ▶ The ps, top, kill and killall commands give you control over running processes.
- ▶ The “ps” command lists processes:

[user@pc]\$ ps	(Current terminal only)
[user@pc]\$ ps -a	(All your processess)
[user@pc]\$ ps -e	(All processes)
- ▶ “top” provides an interactive interface, where you can see the CPU/memory resources taken by each process. (A text-based task manager.)

Process Killing

- ▶ You can send “signals” to processes, often to end them.
- ▶ The “kill” command does this, if you know the process ID:


```
[user@pc]$ kill 1375
```

 (Use ps to find the PID.)
- ▶ By default, kill sends the “SIGTERM” signal.
 - ▶ The process can intercept this to facilitate a graceful exit (e.g. closing files).
- ▶ You can instead send the “SIGKILL” signal, which immediately ends a process (like a “force quit” option):


```
[user@pc]$ kill -KILL 1375
```
- ▶ “killall” lets you give a command name or username:

```
[user@pc]$ killall firefox
```

```
[user@pc]$ killall -u username
```

Daemons

- ▶ Some long-running system processes — called *daemons* — have their very own username and UID (user ID).
- ▶ FTP servers often run as a virtual user called “ftp”.
- ▶ Other servers sometimes run as “nobody”.
- ▶ This helps protect the rest of the system from security vulnerabilities in such software.
- ▶ If the FTP server has a bug, any damage is limited to things the “ftp” user can do.
- ▶ The alternative is for this software to run as “root”, but only if the administrator *really, really* trusts it!

Groups

- ▶ Users can belong to groups.
 - ▶ Each user can belong to any number of groups.
 - ▶ Each group can have any number of users.
- ▶ Normally this is decided by the administrator.
- ▶ Groups allow complex file-ownership/access arrangements.
- ▶ Each group has its own name and a GID (group ID).

File Permissions

- ▶ Each file and directory is owned by both a user and a group.
- ▶ Each also has a set of 9 flags, saying who can do what to it.
- ▶ The flags are “Read”, “Write” and “eXecute”. They are repeated for the user, the group, and everyone else.
 - ▶ Read permission allows you to view a file or directory.
 - ▶ Write permission allows you to modify it.
 - ▶ To delete a file, you need write permission *to its containing directory*, not to the file itself!
 - ▶ Execute permission allows you to run a program or script, or change into a directory.
- ▶ The file’s owner decides:
 - ▶ Their own RWX permissions.
 - ▶ The group’s RWX permissions.
 - ▶ Everyone else’s RWX permissions.

Changing File Permissions

- ▶ File permissions can be changed with the `chmod` command:


```
[user@pc]$ chmod 640 file.txt
```
- ▶ A 3-digit number gives the new permissions: the 1st digit for the owner, the 2nd for the group, and the 3rd everyone else.
- ▶ Each digit is made by adding a combination of:
 - 4 to allow reading,
 - 2 to allow writing, and/or
 - 1 to allow execution.
- ▶ Thus, 640 allows:
 - ▶ reading and writing by the owner ($6 = 4 + 2$);
 - ▶ reading-only by the group (4);
 - ▶ no access to anyone else (0).
- ▶ These are “octal” (base-8) numbers!

Octal File Permissions

- ▶ The complete set of permission combinations are:
 - 0 — no permission.
 - 1 — execute-only.
 - 2 — write-only.
 - 3 — write and execute (2 + 1).
 - 4 — read-only.
 - 5 — read and execute (4 + 1).
 - 6 — read and write (4 + 2).
 - 7 — read, write and execute (4 + 2 + 1).
- ▶ Thus:
 - 600 — only the owner can read and write.
 - 644 — the owner can read/write; everyone else can read.
 - 664 — the owner and group can read/write; everyone else can read.
 - 755 — the owner can do anything; everyone else can read/execute.
 - 777 — unlimited permission (don't use!)

Timestamps

- ▶ All operating systems keep track of the exact time each file was last modified.
- ▶ Some UNIX systems also keep track of:
 - ▶ The original creation time.
 - ▶ The last access time.
- ▶ Any two files can be compared to see which is “newer”.
- ▶ make does this for its targets and pre-requisites (for instance).
- ▶ The touch command can alter a timestamp; by default, setting it to the current time:

```
[user@pc]$ touch program.c
```

(This would cause make to re-compile program.o.)

Symbolic Links (Symlinks)

- ▶ A symlink is a special file that behaves a bit like a pointer — it points to another file (or directory).
- ▶ You can create a symlink with the ln command:


```
[user@pc]$ ln -s originalfile newlink
```

This creates a symlink called “newlink” that references the file “originalfile”.
- ▶ A symlink can be in a completely different directory to the file it references.
- ▶ A symlink takes its permissions from the file it references.
- ▶ You can access/modify the original file by referring to the symlink.

Hard Links

- ▶ A hard link is the connection between a file's data and its directory entry — its name and containing directory.
- ▶ Every file has at least one hard link — one name and location.
- ▶ In some cases, a file can have more than one.
- ▶ A file with two hard links looks like two separate files, which happen to share the same data.
- ▶ This has very specific applications (like incremental backups).
- ▶ However, symlinks are more flexible:
 - ▶ A file with multiple hard links still only has one owner, one group, and one set of permissions.
 - ▶ Hard links can only exist within a single filesystem. Symlinks can link to files on a completely separate drive.

Hard Links and File Deletion

- ▶ On UNIX, deletion is also called “unlinking”.
- ▶ You’re really just removing a hard link.
- ▶ If you remove the *last* hard link, the file is physically deleted.
- ▶ ...except if it’s currently open! If you delete a file while a program is trying to read from or write to it, the file’s directory entry will vanish, but the file itself will remain — anonymously — until closed.

Devices

- ▶ UNIX has other special files that represent hardware devices, or parts of the operating system.
- ▶ These are conventionally located in the `/dev` directory.
- ▶ Some of these are “block” devices that store a block of data.
 - ▶ Hard drives and USB flash drives are block devices.
- ▶ Others are “character” devices that accept or generate a sequence of characters.
 - ▶ The terminal is a character device.
- ▶ Some are virtual devices that provide specialised services (and are useful in scripting).

Common Device Files

Filename	Description
<code>/dev/sda</code>	The first hard drive.
<code>/dev/sda1</code>	The first partition of the first hard drive.
<code>/dev/sdb</code>	The second hard drive (often a USB flash drive).
<code>/dev/cdrom</code>	The CD/DVD drive.
<code>/dev/tty</code>	The current terminal.
<code>/dev/null</code>	A “black hole” — swallows anything you write to it, but is always empty.
<code>/dev/zero</code>	Similar, but contains infinite ‘\0’ bytes.
<code>/dev/urandom</code>	Similar, but contains infinite random bytes.

- ▶ This is a common way to discard the output of a command:

```
[user@pc]$ ./program >/dev/null
```

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Filenames

The filenames — fairly self-explanatory!

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

File Type

- ▶ Here, “backup” is a directory (“d”).
- ▶ The other entries are ordinary files.
- ▶ File types include: “-” for ordinary files, “d” for directories, “l” for symbolic links, “b” for block devices, “c” for character devices, and others for various inter-process communication.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Owner's Permissions

The owner has:

- ▶ Read (“r”) and write (“w”) permission for all entries.
- ▶ Execute (“x”) permission for “backup” and “prog”.
 - ▶ The owner can go into the “backup” directory.
 - ▶ The owner can run the “prog” file.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Group's Permissions

The group has:

- ▶ Read (“r”) permission for all entries.
- ▶ Execute (“x”) permission for “backup” and “prog”.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Everyone Else's Permissions

- ▶ Nobody else can access “backup” or “prog” in any way.
- ▶ Everyone can else read “prog.c” and “prog.h”.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Hard Links

- ▶ All three files have one hard link, which is normal.
- ▶ The backup directory has two. Like all directories, it contains a reference to itself — a special entry called “.”. If it contained subdirectories, its link count would be higher still, due to the subdirectories' special “.” entries.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Owner

All files are owned by the username “12345678”.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Group

- ▶ All files are also owned by the group “student”.
- ▶ Any users belonging to this group will have the group's level of access to these files.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

File Size

The total number of bytes occupied by each file.

Permissions, Owners and Everything Else

The “-l” (for long) switch on `ls` gives a lot of information:

```
[user@pc]$ ls -l
```

```
drwxr-x--- 2 12345678 student 4096 Sep 10 10:07 backup
-rwxr-x--- 1 12345678 student 6719 Sep 10 09:59 prog
-rw-r--r-- 1 12345678 student 799 Sep 10 09:51 prog.c
-rw-r--r-- 1 12345678 student 98 Sep 10 09:52 prog.h
```

Timestamp

The time that each entry was last modified.

Back to Scripting

- ▶ The preceding discussion will prepare you for the environment in which you write scripts.
- ▶ Scripts often need to deal with file permissions, timestamps and process control.

The “if” Statement

```
if command1; then
...
elif command2; then
...
elif command3; then
...
else
...
fi
```

- ▶ The “elif” (else-if) “else” parts are optional.
- ▶ The commands after “if” and “elif” are run to retrieve their exit statuses.
- ▶ These exit statuses act as boolean conditions.
- ▶ “fi” marks the end of an if statement.

Conditions

- ▶ An if statement often looks like this:

```
if [ -e $file ]; then
    ...
fi
```

Technically, “[” is a command by itself.

- ▶ The “[” command is commonly used with if statements.
- ▶ “[” takes a set of parameters that represent a condition to be tested.
 - ▶ e.g. [-e \$file] tests whether the \$file variable refers to an existing file.
- ▶ Like all commands, “[” returns an exit status, which the if statement checks.

File Tests

- ▶ The “[” command can check various file attributes:
 - e thing — does “thing” exist?
 - f thing — is “thing” an ordinary file?
 - d thing — is “thing” a directory?
 - l thing — is “thing” a symlink?
 - r thing — is “thing” readable?
 - w thing — is “thing” writable?
 - x thing — is “thing” executable?
- ▶ Example 1: Check whether “logfile” is writable (by you).

```
if [ -w logfile ]; then ...
```

- ▶ Example 2: Checks whether the value stored in \$dir is a directory.

```
if [ -d ${dir} ]; then ...
```

More Tests

- ▶ The “[” command can perform a few binary tests as well:
 - file1 -nt file2 — is file1 newer than file2?
 - file1 -ot file2 — is file1 older than file2?
 - string1 = string2 — is string1 equal to string2?
 - string1 != string2 — is string1 not equal to string2?
- ▶ And it has separate test for strings containing integers:
 - \$x -eq \$y — is \$x = \$y (numerically)?
 - \$x -ne \$y — is \$x ≠ \$y?
 - \$x -gt \$y — is \$x > \$y?
 - \$x -ge \$y — is \$x ≥ \$y?
 - \$x -lt \$y — is \$x < \$y?
 - \$x -le \$y — is \$x ≤ \$y?
- ▶ When comparing integers, bash ignores leading zeroes and spaces.
- ▶ When comparing strings, bash compares them exactly.

Other Commands as Tests

- ▶ “[” is not the only way.
- ▶ Most commands report “success” or “failure”.
- ▶ Example with grep:

```
if grep -q fail logfile; then
    echo Badness happened\!
fi
```

- ▶ Runs the command “grep -q fail logfile”.
 - ▶ This reports success if it finds “fail” inside logfile.
 - ▶ (The -q silences grep’s terminal output.)
- ▶ Example with ping:

```
if ping -c 1 ${server}; then
    echo I can access ${server}.
fi
```

Boolean Operators

- ▶ The `!`, `&&` and `||` operators (“not”, “and” and “or”) are used in bash as well.
- ▶ They operate on the exit statuses of commands.
- ▶ For instance: “`command1 && command2`” reports success if both commands report success.
- ▶ This is useful in `if` statements as well:

```
if [ -f $file ] && [ $file -nt $file2 ]; then...
```

If `$file` exists, AND it's newer than `$file2`, then run the commands in the `if` statement.

- ▶ Another example:

```
if ! [ -e $file ] || [ $val1 -le 15 ]; then...
```

If `$file` *does not* exist, OR `val1` is less than or equal to 15, then...

Short-Circuit Evaluation

- ▶ In “`cmd1 && cmd2`”, `cmd2` will *only* run if `cmd1` succeeded.
- ▶ In “`cmd1 || cmd2`”, `cmd2` will *only* run if `cmd1` failed.
- ▶ Most programming languages do something similar.
 - ▶ (i.e. they skip the remaining expressions if they already know the overall outcome.)
- ▶ Some script writers use this as a quick-and-dirty `if` statement:

```
[ -f $file ] && cat $file
```

“`cat`” will only run if “`[`” succeeds. This is equivalent to:

```
if [ -f $file ]; then
    cat $file
fi
```

(Note: “`cat $file`” displays the file referred to by `$file`.)

The case Statement

More flexible than Java's `switch` statement:

```
case value in
    (pattern1)
        ...
    ;;
    (pattern2)
        ...
    ;;
    (pattern3)
        ...
    ;;
    ...
esac
```

value is tested against each pattern. If a pattern matches, its commands are run (and we skip the other patterns).

Pattern Matching

Each pattern can contain “`*`”, “`?`”, “`[]`”, “`~`”, etc., with the same meaning as on the command-line (as in Chapter 2). For example:

```
case $thing in
    (abc)
        echo Is abc.
        echo Hurrah.
    ;;
    (abc*) echo Starts with abc. ;;
    (*abc) echo Ends with abc. ;;
    (*abc*) echo Contains abc. ;;
esac
```

- ▶ The value in `$thing` is checked against each pattern in turn.
- ▶ If one matches, the corresponding command(s) will be run.
- ▶ (Note: these are *not* regexes!)

The while Loop

```
while command; do
    ...
done
```

- ▶ Works largely the same way as bash's if statement, except of course as a loop.
- ▶ For example:

```
read filename
while ! [ -e $filename ]; do
    echo $filename does not exist.
    echo Enter another filename:
    read filename
done
```

- ▶ Note: there is no “do-while” loop in bash.

The for (each) Loop

- ▶ Unlike C, bash supports the “for-each”² loop.
 - ▶ It also supports traditional C-style for loops, slightly modified.
 - ▶ However, these are of limited use in a string-based language.
- ▶ For-each loops are simple: you have a list of values, and the loop gives you one at a time.
- ▶ On each iteration, the loop picks the next value in the list, and sets a variable to that value.
- ▶ The commands in the loop then deal with that variable.

```
for variable in list-of-values; do
    ...
done
```

²The word “each” doesn't actually appear in bash — that's just a generic term.

Input

- ▶ You can acquire user input with the read command, which reads a line of text:

```
read variable
```

- ▶ You can also read multiple values at once:

```
read variable1 variable2 variable3
```

- ▶ The user is expected to enter at least three words.
- ▶ The 1st will be assigned to **variable1**, the 2nd to **variable2** and the rest of the line to **variable3**.

Command-line parameters

- ▶ Script files can have their own command-line parameters.
- ▶ These are represented by special variables: \$1 for the first parameter, \$2 for the second, and so on.
- ▶ Say you create a file called thescript, containing this:

```
#!/bin/bash
echo $1 is $2 years old and lives at $3.
```

- ▶ You could run this as follows:

```
[user@pc]$ ./thescript Barry 102 "14 Grass St."
```

Based on this, \$1 would be “Barry”, \$2 would be “102” and \$3 would be “14 Grass St.”.

- ▶ There are other special variables; e.g.
 - ▶ \$# is the number of parameters (3, in the above example).
 - ▶ \$* combines all parameters (“Barry 102 14 Mountain St.”).

Arithmetic

- ▶ bash supports rudimentary maths expressions (integer-only).
- ▶ To evaluate an expression, you can enclose it in `$((and))`.
- ▶ This works by substitution — the whole expression is *replaced* by its result.
- ▶ For instance: `$((5 + 6 * 4))` is evaluated and replaced by `"29"`.
- ▶ You can embed variables in these expressions as well.
- ▶ You can use the result (practically) anywhere.

Command Substitution

- ▶ You can form a command from the output of other commands.
- ▶ Enclose a command in `$(...)`, or ``...`` (backticks):
 - ▶ Don't confuse the 1st with `$(...)` for arithmetic.
 - ▶ Don't confuse the 2nd with single quotes, `'...'`.
- ▶ Bash runs the command and captures its output, which then replaces the command.
- ▶ For example:

```
thedata="$(date)"
```

- ▶ Run by itself, `"date"` shows the current date and time.
- ▶ Here, this is output is captured and assigned to the variable `thedata`.

Command Substitution — Examples

```
for word in $(cat file.txt); do
    echo $word
done
```

- ▶ `$(cat file.txt)` gets the entire contents of `file.txt`.
- ▶ This causes the `for` loop to iterate over each word in the file.
- ▶ Thus, each word in `file.txt` is printed on a separate line.

```
if [ $(date +%A) = Sunday ]; then
    echo "Free ice cream for everyone!"
fi
```

If the current day (as determined by `"date +%A"`) is Sunday, output `"Free ice cream for everyone!"`

Exit Statuses — True or False

- ▶ Most decision making in Bash depends on "exit statuses".
- ▶ When a program/command ends, it returns a single integer to the operating system — its exit status.
- ▶ Zero indicates success. Anything else indicates failure.
 - ▶ (This is the opposite way around from boolean values in C!)
- ▶ We've been using exit statuses all along:

```
int main(void) {
    ...
    return 0;    /* Exit status == success! */
}
```

- ▶ We *could* instead return a variable, indicating either success or failure. Many programs do this.

C and Bash Interacting

myprogram.c

```
int main(void) {
    int status = 1; /* "Failure" by default */
    ...
    if(...) {        /* Error checking */
        ...          /* Do useful things... */
        status = 0; /* "Success"! */
    }
    return status;
}
```

Then in Bash:

```
if ./myprogram; then # Run the above C program,
    echo "It worked!" # and check if it succeeded.
fi
```

47/51

Regular Expressions, or “Regex”es

- ▶ We’ve already seen Bash’s pattern matching.
- ▶ Regexes are the same idea, but with different notation.
- ▶ Used by various UNIX commands; e.g. grep, sed, awk, vim.
- ▶ *Wildcards* stand for a single unknown character:
 - ▶ “.” — any character.
 - ▶ “[abcm-z]” — any character within the brackets, where “-” gives a range of characters.
- ▶ *Quantifiers* may appear after something else, and say how many times it should occur:
 - ▶ “?” — zero or one times.
 - ▶ “*” — zero or more times.
 - ▶ “+” — one or more times.
- ▶ Thus:
 - ▶ “.*” means any sequence of characters.
 - ▶ “abc?” means “ab”, optionally followed by “c”.
 - ▶ “[0-9]+” means any non-empty digit sequence.

48/51

grep

- ▶ With grep, you can find patterns in files, streams and throughout entire directory hierarchies.
 - ▶ Use -E to allow “extended” regular expressions.
 - ▶ -E is not always required, but we’ll use it for simplicity.

- ▶ Example 1: find all integers in file.txt:

```
grep -E '[0-9]+' file.txt
```

- ▶ Example 2: find all words ending in “ism” in the output of somecprogram:

```
./somecprogram | grep -E '[a-zA-Z]*ism'
```

- ▶ Example 3: find all occurrences of fprintf and sprintf throughout all files in the code directory:

```
grep -E -R '[fs]printf\(' code/
```

49/51

More Regex Notation

- ▶ Prefix a special character with “\” to take it literally:
 - ▶ “\.*” means “.*”.
 - ▶ “\?+” means a sequence of one or more “?”s.
 - ▶ “\?” means an optional “\”.
- ▶ You can group characters with “(...)”:
 - ▶ “a(bc)?” means “a” optionally followed by “bc”.
 - ▶ “(silly)*me” means any number of “silly ”s, followed by “me”; e.g. “silly silly silly me”.
- ▶ You can select between alternatives with “|”:
 - ▶ “abc|def|ghi” means either “abc” or “def” or “ghi”.
 - ▶ “(heads |tails)+” means one or more of either “heads” or “tails”; e.g. “heads heads tails heads tails”.
- ▶ You can refer to the beginning/end of the line:
 - ▶ “^abc” means “abc”, but only at the beginning of the line.
 - ▶ “abc\$” means “abc”, but only at the end of the line.

50/51

vim, less and man

- ▶ The `vim` editor and the `less` viewer both support regex searching.
- ▶ Type `/`, enter a regex, and press Enter.
 - ▶ When a match is found, you can find the next one by typing `/` and Enter again.
- ▶ The `man` command usually uses `less`, so you can do regex searching while viewing man pages.
 - ▶ Try it with the bash man page!